

# DEEP LEARNING

## HANDOUT: INTRODUCTION TO TENSORFLOW

### 1. Introduction to TensorFlow

TensorFlow is an open-source deep learning and machine learning framework developed by Google. It provides a comprehensive, flexible ecosystem of tools, libraries, and community resources for building and deploying machine learning applications across a wide range of platforms.

### 2. Role of TensorFlow in Building Deep Learning Models

- TensorFlow simplifies the model-building workflow across all stages:
- Data Ingestion: Load and preprocess data using `tf.data` pipelines.
- Model Definition: Use `tf.keras.Sequential` or the Functional API to define neural networks.
- Training: Compile models with optimizers and loss functions and fit them using `.fit()`.
- Evaluation and Prediction: Use `.evaluate()` and `.predict()` on unseen data.
- Deployment: Export models with `SavedModel` or convert to mobile formats with TensorFlow Lite.

### 3. Key Features:

- Cross-Platform: Runs on CPUs, GPUs, TPUs, and mobile devices.
- Scalability: Suitable for training models on both local machines and distributed servers.
- Keras Integration: Offers a user-friendly API for rapid model building.
- Auto Differentiation: Simplifies backpropagation using `tf.GradientTape`.

### 4. What Are Tensors?

Tensors are multi-dimensional arrays used to represent data in deep learning models. They are the building blocks of all computations in TensorFlow — storing inputs, outputs, weights, gradients, etc. Tensors generalize matrices to higher dimensions. For example:

Scalar: 0-D tensor (e.g., 5)

Vector: 1-D tensor (e.g., [5, 6])

Matrix: 2-D tensor (e.g., [[1, 2], [3, 4]])

3D and beyond: for image, video, and sequence data

## 5. How Tensors Are Used to Build Neural Networks?

In TensorFlow, all model inputs, outputs, and intermediate data are represented as tensors. These tensors flow through the layers of a model, undergoing transformations through tensor operations like matrix multiplication, addition, and nonlinear activation functions.

For example, consider a simple dense neural network layer. The input tensor is multiplied by a weight tensor and then passed through an activation function to produce an output tensor:

```
import tensorflow as tf  
  
x = tf.constant([[1.0, 2.0]]) # input tensor (1x2)  
  
layer = tf.keras.layers.Dense(1) # single dense layer  
  
output = layer(x) # output tensor (1x1)  
  
print(output)
```

### Code Explanation:

- x is a 2D tensor with shape [1, 2]
- The layer has a weight tensor of shape [2, 1]
- tf.matmul(x, weights) + bias is computed internally
- Result is passed through an activation (default: linear)

*Tensors are not just containers for data; they enable efficient computation, automatic differentiation, and hardware acceleration on GPUs and TPUs.*

### Why does TensorFlow emphasize tensors instead of arrays?

**Answer:** Tensors support automatic differentiation, GPU execution, and dynamic shape manipulation.

**Practice Task:** Build a simple dense neural layer using `tf.keras.layers.Dense` and pass a 2D tensor input through it.

```
x = tf.constant([[1.0, 2.0], [3.0, 4.0]])
dense = tf.keras.layers.Dense(3)
output = dense(x)
print(output)
```

### Code Explanation:

#### Line 1: Creating Input Tensor

```
x = tf.constant([[1.0, 2.0], [3.0, 4.0]])
```

- This creates a 2D tensor (or matrix) with shape (2, 2):

[[1.0, 2.0],

[3.0, 4.0]]

- Each row is treated as a separate input sample, and each column is a feature.
- So we are feeding 2 samples, each with 2 features, into the neural network.

#### Line 2: Defining a Dense Layer

```
dense = tf.keras.layers.Dense(3)
```

- Creates a Dense (fully connected) layer with:
  - 3 units (neurons) in the output
  - Input shape inferred from the input x (which has 2 features per sample)
  - Weights (W) shape: (2, 3)
  - Bias (b) shape: (3,)
- The output of this layer will be a tensor of shape (2, 3):
  - 2 rows for 2 input samples
  - 3 values per row for 3 neurons

#### Line 3: Forward Pass (Layer Execution)

```
output = dense(x)
```

This line performs the following operation for each input row:

output=xW+b

Where:

- x: input tensor of shape (2, 2)
- W: weights of shape (2, 3)
- b: bias of shape (3,)
- Result: tensor of shape (2, 3)

Example pseudo computation (not actual values):

```
output[0] = [1*W11 + 2*W21 + b1, 1*W12 + 2*W22 + b2, 1*W13 + 2*W23 + b3]
```

```
output[1] = [3*W11 + 4*W21 + b1, 3*W12 + 4*W22 + b2, 3*W13 + 4*W23 + b3]
```

**Note: The weights and biases are initialized randomly unless manually set.**

#### Line 4: Output

```
print(output)
```

This will print a 2x3 matrix (a tensor of shape (2, 3))

```
tf.Tensor(  
[[ 0.123 -0.456 0.789]  
 [ 0.321 -0.654 0.987]], shape=(2, 3), dtype=float32)
```

Each row corresponds to the dense layer output for one input sample.

#### Summary Table

Component	Shape	Meaning
Input x	(2, 2)	2 samples, 2 features each
Weights W	(2, 3)	Learnable parameters
Bias b	(3,)	Learnable offsets for each neuron
Output	(2, 3)	Dense layer result for 2 samples

#### 6. Data Types in TensorFlow

TensorFlow supports multiple data types (dtypes) that affect memory use, precision, and speed. Choosing the right type ensures efficiency and accuracy.

Category	Examples	Description
Floating-point	tf.float16, tf.float32, tf.float64	Decimal numbers
Integer	tf.int8, tf.int16, tf.int32	Whole numbers
Boolean	tf.bool	True/False values
String	tf.string	Text data
Complex	tf.complex64, tf.complex128	Complex numbers

### Example

```
import tensorflow as tf

a = tf.constant(3.14, dtype=tf.float32)

b = tf.constant(42, dtype=tf.int64)

c = tf.constant("Hello", dtype=tf.string)
```

Your model needs to process both images (0-255 integers) and text data. Which dtypes would you use?

**Answer:** Images: tf.uint8, Text: tf.string

You encounter a type mismatch error between operations. What will you check?

**Answer:** Ensure all tensors in the operation share compatible dtypes.

**Problem:** Create a tensor of shape [4, 4] with float values, then cast it to int32.

**Solution:**

```
import tensorflow as tf

float_tensor = tf.random.uniform([4, 4], dtype=tf.float32)

int_tensor = tf.cast(float_tensor, tf.int32)

print("Original tensor:", float_tensor)

print("Casted tensor:", int_tensor)
```

## 7. Numerical Precision

Precision controls how numbers are stored:

- float16 is faster but may lose accuracy (underflows)
- float32 is standard for neural networks
- float64 is slower but preserves precision

Too low precision can lead to numerical instability.

### Example

```
tiny_float16 = tf.constant(1e-7, dtype=tf.float16)
```

```
tiny_float32 = tf.constant(1e-7, dtype=tf.float32)
```

Your model outputs NaN during training. What are two possible fixes?

**Answer:** Use `tf.clip_by_norm()` or switch to float32.

You observe inconsistent outputs across devices (CPU vs GPU). Could precision be a factor?

**Answer:** Yes. Different hardware may use different default precisions.

### Practice Problem

Try computing the dot product of two large vectors using float16 and float64. Observe differences.

#### Solution:

```
import tensorflow as tf

v1 = tf.random.uniform([1000000], dtype=tf.float16)
v2 = tf.random.uniform([1000000], dtype=tf.float16)
dot_16 = tf.tensordot(v1, v2, axes=1)

v1_64 = tf.cast(v1, tf.float64)
v2_64 = tf.cast(v2, tf.float64)

dot_64 = tf.tensordot(v1_64, v2_64, axes=1)

print("Dot product with float16:", dot_16.numpy())
print("Dot product with float64:", dot_64.numpy())
```

## 8. Tensors vs Variables

Tensors in TensorFlow are immutable. To store weights or other learnable parameters that change during training, use `tf.Variable`, which allows updates.

Feature	<code>tf.Tensor</code>	<code>tf.Variable</code>
Mutability	Immutable	Mutable
Use Case	Input data	Trainable model weights

### Example

```
weights = tf.Variable(tf.random.normal([3, 2]))
```

```
weights.assign_add(tf.ones([3, 2]))
```

**Why can't we use regular tensors for neural network weights?**

**Answer:** Because weights need to be updated during training.

**How would you freeze a pretrained layer in TensorFlow?**

**Answer:** Set `trainable=False` on its variables.

### Practice Problem

Create a trainable variable initialized with zeros of shape [5,5] and increment its values by 1.

#### Solution:

```
import tensorflow as tf  
  
var = tf.Variable(tf.zeros([5, 5]))  
  
var.assign_add(tf.ones([5, 5]))  
  
print(var)
```

## 9. Creating Tensors

TensorFlow provides a variety of methods to create tensors for different purposes from constant values to random initialization. These are fundamental when building models.

Method	Example	Use Case
<code>tf.constant()</code>	<code>tf.constant([[1, 2], [3, 4]])</code>	Fixed data
<code>tf.zeros()</code>	<code>tf.zeros([2, 3])</code>	Initialization

<code>tf.ones()</code>	<code>tf.ones([2, 2])</code>	Fill with ones
<code>tf.eye()</code>	<code>tf.eye(3)</code>	Identity matrix
<code>tf.random.normal()</code>	<code>tf.random.normal([2,2], mean=0, stddev=1)</code>	Random weights

How would you initialize a weight matrix with small random values for a neural layer?

**Answer:** Use `tf.random.normal()` with small stddev.

You want to create a batch of 5 grayscale images (28x28 pixels) initialized to 0. What method would you use?

**Answer:** `tf.zeros([5, 28, 28])`

### Practice Problem

Create a 3x3 tensor with values ranging from 0 to 8.

#### Solution:

```
import tensorflow as tf
t = tf.reshape(tf.range(9), [3, 3])
print(t)
```

### 10. Indexing and Slicing

TensorFlow supports powerful slicing and indexing operations similar to NumPy. These help extract or manipulate subsets of data.

Type	Example	Description
Basic	<code>tensor[0]</code>	First element or row
Slicing	<code>tensor[1:3, :2]</code>	Slice rows and columns
Boolean	<code>tensor[tensor &gt; 0]</code>	Masked selection

#### Example

```
t = tf.constant([[1, 2, 3], [4, 5, 6]])
print(t[:, 1:]) # [[2, 3], [5, 6]]
```

How do you access the second column of a matrix?

**Answer:** tensor[:, 1]

Extract every even-indexed row from a batch of 10 images.

**Answer:** tensor[::2]

### Practice Problem

Given a 4x4 tensor, extract the diagonal elements.

**Solution:**

```
t = tf.reshape(tf.range(16), [4, 4])  
diagonal = tf.linalg.diag_part(t)  
print(diagonal)
```

## 11. Dimensional Operations

These operations allow changing or manipulating tensor shapes to match the requirements of models or functions.

Operation	Example	Effect
reshape()	tf.reshape(t, [6,10])	Changes shape
expand_dims()	tf.expand_dims(t, axis=0)	Adds new dimension
squeeze()	tf.squeeze(t)	Removes dim=1

How to convert a batch of images from [28,28,1] to [28,28]?

**Answer:** tf.squeeze(t, axis=-1)

### Practice Problem

Create a tensor of shape [3, 1, 4] and convert it to [3, 4].

**Solution:**

```
t = tf.random.normal([3, 1, 4])
```

```
t_flat = tf.squeeze(t)  
print(t_flat.shape)
```

## 12. Broadcasting

Broadcasting lets TensorFlow automatically expand tensors of smaller shapes to match larger ones when performing element-wise operations.

### Rules

- Align shapes from the right
- If dimensions differ and one of them is 1, broadcasting occurs

### Example

```
a = tf.constant([1, 2, 3])    # (3,)  
  
b = tf.constant([[10], [20]]) # (2,1)  
  
print(a + b)              # [[11,12,13],[21,22,23]]
```

What's the shape of [3,1] \* [1,4]?

Answer: [3,4]

You get a shape mismatch error in element-wise multiplication. What's likely the issue?

Answer: Incompatible shapes without broadcasting rules being satisfied.

**Practice Problem:** Multiply a tensor of shape [2,3] with [3] using broadcasting.

### Solution:

```
a = tf.constant([[1, 2, 3], [4, 5, 6]])  
  
b = tf.constant([1, 2, 3])  
  
print(a * b)
```

## 13. Math Operations

TensorFlow supports a wide range of mathematical operations, including element-wise arithmetic, matrix multiplication, and reductions like mean or sum.

Type	Examples	Description
------	----------	-------------

Arithmetic	<code>tf.add()</code> , <code>tf.subtract()</code>	Element-wise operations
Matrix	<code>tf.matmul()</code> , <code>tf.linalg.inv()</code>	Matrix algebra
Reduction	<code>tf.reduce_sum()</code> , <code>tf.reduce_mean()</code>	Aggregation functions

### Example

```
a = tf.constant([[1, 2], [3, 4]])

print(tf.reduce_sum(a, axis=1)) # Output: [3, 7]
```

How do you compute the mean of all elements in a tensor?

**Answer:** `tf.reduce_mean(tensor)`

What is the difference between `tf.multiply()` and `tf.matmul()`?

**Answer:** `tf.multiply()` is element-wise; `tf.matmul()` is matrix multiplication.

**Practice Problem:** Compute the dot product of two vectors  $[1, 2, 3]$  and  $[4, 5, 6]$ .

### Solution:

```
v1 = tf.constant([1, 2, 3])
v2 = tf.constant([4, 5, 6])
dot_product = tf.tensordot(v1, v2, axes=1)
print(dot_product) # Output: 32
```

## 14. Merge and Split

TensorFlow allows joining and separating tensors for reshaping or batching operations.

Operation	Example	Use Case
<code>concat()</code>	<code>tf.concat([a, b], axis=0)</code>	Join along an axis
<code>stack()</code>	<code>tf.stack([a, b], axis=0)</code>	Create new dimension
<code>split()</code>	<code>tf.split(tensor, num=3, axis=1)</code>	Split into parts

How would you combine two tensors of shape [2,3] vertically?

**Answer:** Use `tf.concat([a, b], axis=0)`

You need to split a 6x2 tensor into three 2x2 matrices. Which function do you use?

**Answer:** `tf.split(tensor, num_or_size_splits=3, axis=0)`

**Practice Problem:** Stack two 1D tensors [1, 2, 3] and [4, 5, 6] to make a 2x3 tensor.

**Solution:**

```
a = tf.constant([1, 2, 3])  
b = tf.constant([4, 5, 6])  
st = tf.stack([a, b], axis=0)  
print(st)
```

## 15. Statistics

Statistical operations help summarize or rank values inside tensors.

Function	Example	Description
<code>reduce_mean()</code>	<code>tf.reduce_mean(t)</code>	Mean of all elements
<code>argmax()</code>	<code>tf.argmax(t, axis=1)</code>	Index of max values
<code>top_k()</code>	<code>tf.math.top_k(t, k=3)</code>	Top K values and indices

How do you find the most likely class from softmax output?

**Answer:** Use `tf.argmax(predictions, axis=1)`

You want to find the 3 highest predictions from a tensor. What function do you use?

**Answer:** `tf.math.top_k()`

**Practice Problem:** Given `t = [3.2, 1.5, 7.8, 5.5]`, return the top 2 values.

**Solution:**

```
t = tf.constant([3.2, 1.5, 7.8, 5.5])  
values, indices = tf.math.top_k(t, k=2)  
print("Top values:", values.numpy())
```

## 16. Tensor Comparison

TensorFlow supports comparison operations such as equality, greater-than, and conditional selection.

Operation	Example	Description
tf.equal()	tf.equal(a, b)	Element-wise equality
tf.greater()	tf.greater(a, b)	Element-wise greater check
tf.where()	tf.where(condition, x, y)	Element from x or y based on condition

How do you replace all negative values in a tensor with 0?

**Answer:** tf.where(tensor < 0, 0, tensor)

How would you find where two tensors have equal values?

**Answer:** Use tf.equal(tensor1, tensor2)

**Practice Problem:** Given a tensor t = [-1, 3, -5, 2], replace negative values with 0.

**Solution:**

```
t = tf.constant([-1, 3, -5, 2])
result = tf.where(t < 0, tf.zeros_like(t), t)
print(result)
```

## 17. Fill and Copy

Useful for creating fixed-value tensors or duplicating tensor contents.

Function	Example	Description
tf.fill()	tf.fill([2, 2], 5)	Tensor filled with value
tf.identity()	tf.identity(t)	Clone of input tensor

How do you create a 10x10 tensor filled with value 7?

**Answer:** `tf.fill([10, 10], 7)`

What's the purpose of `tf.identity()` in computation graphs?

**Answer:** It allows keeping computation paths separate and preserved

**Practice Problem:** Create a 3x3 tensor filled with the value -2.

**Solution:**

```
t = tf.fill([3, 3], -2)  
print(t)
```

## 18. Data Limiting

TensorFlow provides utilities to clip tensor values or norms to maintain numerical stability, especially in deep learning.

Method	Use Case
<code>clip_by_value()</code>	Restrict range of values
<code>clip_by_norm()</code>	Prevent exploding gradients

How would you limit pixel intensities to [0, 255]?

**Answer:** `tf.clip_by_value(image, 0, 255)`

**Practice Problem:** Clip values in `t = [-10, 0, 5, 20]` between 0 and 10.

**Solution:**

```
t = tf.constant([-10, 0, 5, 20], dtype=tf.float32)  
clipped = tf.clip_by_value(t, 0, 10)  
print(clipped)
```

## 19. Advanced Operations

TensorFlow includes advanced tensor types and tools for custom training workflows.

Feature	Use Case
GradientTape	Manual gradient computation
SparseTensor	Efficient for sparse data storage
RaggedTensor	Handling sequences of varying lengths

You want to compute gradient of  $y = x^2$  manually. What do you use?

**Answer:** tf.GradientTape()

**Practice Problem:** Use GradientTape to compute the derivative of  $y = x^3$  at  $x = 2$ .

**Solution:**

```
x = tf.Variable(2.0)
with tf.GradientTape() as tape:
    y = x ** 3
    grad = tape.gradient(y, x)
    print(grad)
# Output: 12.0
```